

Alam_intro_to_stat_analysis_final_proj

December 8, 2020

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
[2]: weather = pd.read_csv("C:/Users/shams/OneDrive/Desktop/All_project/
↳Intro_to_stat_analysis_proj/weatherHistory.csv")
weather.head()
```

```
[2]:
```

	Formatted Date	Summary	Precip Type	Temperature (C)	\
0	2006-04-01 00:00:00.000 +0200	Partly Cloudy	rain	9.472222	
1	2006-04-01 01:00:00.000 +0200	Partly Cloudy	rain	9.355556	
2	2006-04-01 02:00:00.000 +0200	Mostly Cloudy	rain	9.377778	
3	2006-04-01 03:00:00.000 +0200	Partly Cloudy	rain	8.288889	
4	2006-04-01 04:00:00.000 +0200	Mostly Cloudy	rain	8.755556	

	Apparent Temperature (C)	Humidity	Wind Speed (km/h)	\
0	7.388889	0.89	14.1197	
1	7.227778	0.86	14.2646	
2	9.377778	0.89	3.9284	
3	5.944444	0.83	14.1036	
4	6.977778	0.83	11.0446	

	Wind Bearing (degrees)	Visibility (km)	Loud Cover	Pressure (millibars)	\
0	251.0	15.8263	0.0	1015.13	
1	259.0	15.8263	0.0	1015.63	
2	204.0	14.9569	0.0	1015.94	
3	269.0	15.8263	0.0	1016.41	
4	259.0	15.8263	0.0	1016.51	

Daily Summary

0	Partly cloudy throughout the day.
1	Partly cloudy throughout the day.
2	Partly cloudy throughout the day.
3	Partly cloudy throughout the day.
4	Partly cloudy throughout the day.

```
[ ]: """The CSV file includes a hourly/daily summary for Szeged, Hungary area,
↳between 2006 and 2016.It is a very large data set
with dimension (96453,12).
My analysis was to find is there a relationship between humidity and
↳temperature? What about between humidity and apparent
Temperature."""
```

```
[ ]: """Data Background"""
"""
4-year air temperature data from June 2014 to May 2018 in 10-min averages were
↳used based on 1-min measurements.
To compare temperature modifying effects of different LCZs in Szeged LCZ
↳averages were used. Consequently,
in case of LCZ 2 and 3 it means only one station, however in LCZ 6 the average
↳of ten stations was used according
to the size of this LCZ class.
Available from:
https://www.researchgate.net/publication/
327867323_Weather_and_climate_modeling_possibilities_using_local_climate_zone_concept_and_
observation_network_in_Szeged_Hungary [accessed Dec 08 2020].
"""
```

```
[26]: weather[weather.isnull().any(axis=1)]
```

```
[26]:
```

	Formatted Date	Summary	Precip	Type	\
52672	2012-04-11 16:00:00.000 +0200	Mostly Cloudy		NaN	
52674	2012-04-11 18:00:00.000 +0200	Mostly Cloudy		NaN	
52675	2012-04-11 19:00:00.000 +0200	Mostly Cloudy		NaN	
52677	2012-04-11 21:00:00.000 +0200	Mostly Cloudy		NaN	
52678	2012-04-11 22:00:00.000 +0200	Mostly Cloudy		NaN	
...		
95584	2016-10-31 19:00:00.000 +0100	Mostly Cloudy		NaN	
95585	2016-10-31 20:00:00.000 +0100	Mostly Cloudy		NaN	
95586	2016-10-31 21:00:00.000 +0100	Mostly Cloudy		NaN	
95587	2016-10-31 22:00:00.000 +0100	Partly Cloudy		NaN	
95588	2016-10-31 23:00:00.000 +0100	Mostly Cloudy		NaN	

	Temperature (C)	Apparent Temperature (C)	Humidity	Wind Speed (km/h)	\
52672	19.016667	19.016667	0.26	14.8764	
52674	17.850000	17.850000	0.28	13.7977	
52675	16.322222	16.322222	0.32	10.8192	
52677	12.566667	12.566667	0.43	9.0160	
52678	12.927778	12.927778	0.47	17.6295	
...	
95584	8.322222	7.044444	0.85	8.0339	
95585	7.627778	6.183333	0.87	8.2271	
95586	7.111111	5.511111	0.89	8.5008	

95587	6.672222	4.961111	0.90	8.6457
95588	6.322222	4.588889	0.91	8.4686


	Wind Bearing (degrees)	Visibility (km)	Loud Cover \
52672	163.0	9.982	0.0
52674	169.0	9.982	0.0
52675	151.0	9.982	0.0
52677	159.0	9.982	0.0
52678	197.0	16.100	0.0
...
95584	290.0	0.000	0.0
95585	293.0	0.000	0.0
95586	297.0	0.000	0.0
95587	299.0	0.000	0.0
95588	299.0	0.000	0.0

	Pressure (millibars)	Daily Summary
52672	1002.40	Mostly cloudy until night.
52674	1001.79	Mostly cloudy until night.
52675	1001.60	Mostly cloudy until night.
52677	1001.92	Mostly cloudy until night.
52678	1002.20	Mostly cloudy until night.
...
95584	1021.73	Mostly cloudy starting in the afternoon.
95585	1021.76	Mostly cloudy starting in the afternoon.
95586	1021.81	Mostly cloudy starting in the afternoon.
95587	1021.83	Mostly cloudy starting in the afternoon.
95588	1021.80	Mostly cloudy starting in the afternoon.

[517 rows x 12 columns]

```
[33]: import seaborn as sns
plt.figure(figsize=(12,8))
weather.isna().sum()
sns.heatmap(weather.isna())
```

[33]: <matplotlib.axes._subplots.AxesSubplot at 0x2823da74688>



```
[5]: """We can see that the Precip Type variable have some missing values, as the
    ↳ dataset is huge we will remove the rows"""
    """Dropping the missing values"""
    dat1 = weather.dropna()
    dat1.isna().sum()
```

```
[5]: Formatted Date      0
    Summary              0
    Precip Type          0
    Temperature (C)      0
    Apparent Temperature (C) 0
    Humidity             0
    Wind Speed (km/h)    0
    Wind Bearing (degrees) 0
    Visibility (km)      0
    Loud Cover           0
```

```
Pressure (millibars)      0
Daily Summary             0
dtype: int64
```

```
[ ]: """We will perform a time analysis on the dataset, as time and season is an
    ↪important componed for weather analysis"""
    """Time Analysis"""
```

```
[6]: dat1['Formatted Date'] = pd.to_datetime(dat1['Formatted Date'], utc=True)
    dat1['Formatted Date']
```

C:\Users\shams\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 """Entry point for launching an IPython kernel.

```
[6]: 0      2006-03-31 22:00:00+00:00
    1      2006-03-31 23:00:00+00:00
    2      2006-04-01 00:00:00+00:00
    3      2006-04-01 01:00:00+00:00
    4      2006-04-01 02:00:00+00:00
    ...
    96448   2016-09-09 17:00:00+00:00
    96449   2016-09-09 18:00:00+00:00
    96450   2016-09-09 19:00:00+00:00
    96451   2016-09-09 20:00:00+00:00
    96452   2016-09-09 21:00:00+00:00
    Name: Formatted Date, Length: 95936, dtype: datetime64[ns, UTC]
```

```
[7]: #Set index as "Date"

    dat_index = dat1.set_index('Formatted Date')
    dat_index.head()
```

```
[7]:
```

Formatted Date	Summary	Precip	Type	Temperature (C)	\
2006-03-31 22:00:00+00:00	Partly Cloudy		rain	9.472222	
2006-03-31 23:00:00+00:00	Partly Cloudy		rain	9.355556	
2006-04-01 00:00:00+00:00	Mostly Cloudy		rain	9.377778	
2006-04-01 01:00:00+00:00	Partly Cloudy		rain	8.288889	
2006-04-01 02:00:00+00:00	Mostly Cloudy		rain	8.755556	

Apparent Temperature (C)	Humidity	\
--------------------------	----------	---

Formatted Date		
2006-03-31 22:00:00+00:00	7.388889	0.89
2006-03-31 23:00:00+00:00	7.227778	0.86
2006-04-01 00:00:00+00:00	9.377778	0.89
2006-04-01 01:00:00+00:00	5.944444	0.83
2006-04-01 02:00:00+00:00	6.977778	0.83

	Wind Speed (km/h)	Wind Bearing (degrees) \
Formatted Date		
2006-03-31 22:00:00+00:00	14.1197	251.0
2006-03-31 23:00:00+00:00	14.2646	259.0
2006-04-01 00:00:00+00:00	3.9284	204.0
2006-04-01 01:00:00+00:00	14.1036	269.0
2006-04-01 02:00:00+00:00	11.0446	259.0

	Visibility (km)	Loud Cover	Pressure (millibars) \
Formatted Date			
2006-03-31 22:00:00+00:00	15.8263	0.0	1015.13
2006-03-31 23:00:00+00:00	15.8263	0.0	1015.63
2006-04-01 00:00:00+00:00	14.9569	0.0	1015.94
2006-04-01 01:00:00+00:00	15.8263	0.0	1016.41
2006-04-01 02:00:00+00:00	15.8263	0.0	1016.51

	Daily Summary
Formatted Date	
2006-03-31 22:00:00+00:00	Partly cloudy throughout the day.
2006-03-31 23:00:00+00:00	Partly cloudy throughout the day.
2006-04-01 00:00:00+00:00	Partly cloudy throughout the day.
2006-04-01 01:00:00+00:00	Partly cloudy throughout the day.
2006-04-01 02:00:00+00:00	Partly cloudy throughout the day.

```
[8]: data_columns = ['Apparent Temperature (C)', 'Humidity']
df_monthly_mean = dat_index[data_columns].resample('MS').mean()
df_monthly_mean.head()
```

	Apparent Temperature (C)	Humidity
Formatted Date		
2005-12-01 00:00:00+00:00	-4.050000	0.890000
2006-01-01 00:00:00+00:00	-4.173708	0.834610
2006-02-01 00:00:00+00:00	-2.990716	0.843467
2006-03-01 00:00:00+00:00	1.969780	0.778737
2006-04-01 00:00:00+00:00	12.098827	0.728625

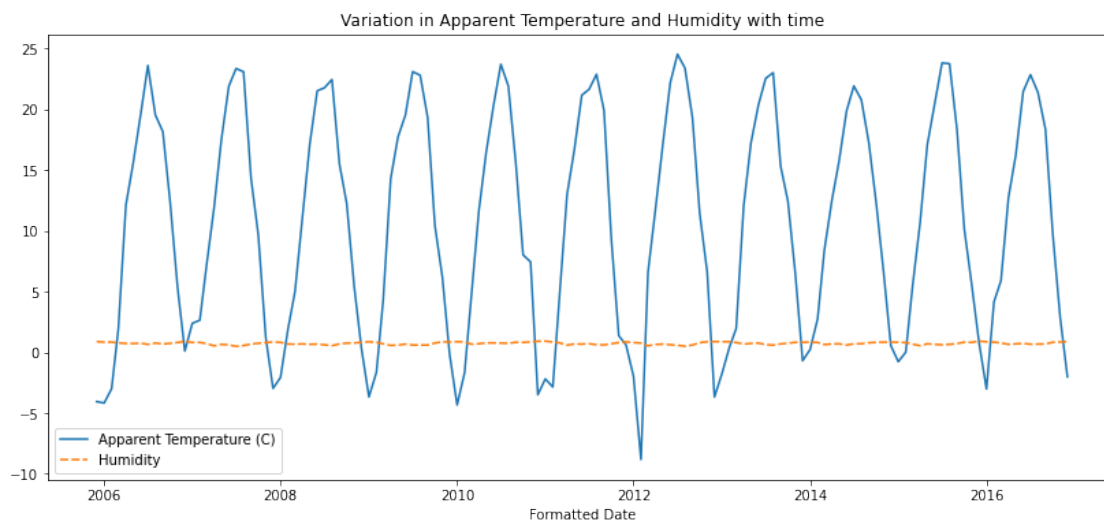
```
[9]: #Plotting Variation in Apparent Temperature and Humidity with time

import seaborn as sns
import warnings
```

```
warnings.filterwarnings("ignore")

plt.figure(figsize=(14,6))
plt.title("Variation in Apparent Temperature and Humidity with time")
sns.lineplot(data=df_monthly_mean)
```

[9]: <matplotlib.axes._subplots.AxesSubplot at 0x28238300f88>



[10]: *"""We can see a sinusoidal shaped variation of the apparent temperature with time which denotes the cycle of season.
From 2006-2016 the Humidity has been pretty constant with no big variation. """*

#retrieving the data of a particular month from every year, say April

```
df1 = df_monthly_mean[df_monthly_mean.index.month==4]
print(df1)
```

```
df1.dtypes
```

Formatted Date	Apparent Temperature (C)	Humidity
2006-04-01 00:00:00+00:00	12.098827	0.728625
2007-04-01 00:00:00+00:00	11.894421	0.536361
2008-04-01 00:00:00+00:00	11.183688	0.693194
2009-04-01 00:00:00+00:00	14.267076	0.567847
2010-04-01 00:00:00+00:00	11.639406	0.706875
2011-04-01 00:00:00+00:00	12.978997	0.591625
2012-04-01 00:00:00+00:00	11.780703	0.643583
2013-04-01 00:00:00+00:00	12.045563	0.677667
2014-04-01 00:00:00+00:00	12.486181	0.691403

2015-04-01 00:00:00+00:00	10.632801	0.547764
2016-04-01 00:00:00+00:00	12.731427	0.659972


```
[10]: Apparent Temperature (C)    float64
      Humidity                    float64
      dtype: object
```

```
[11]: #Plotting each years Humidity and Temperature change
```

```
import matplotlib.dates as mdates
from datetime import datetime

fig, ax = plt.subplots(figsize=(15,5))
ax.plot(df1.loc['2006-04-01':'2016-04-01', 'Apparent Temperature (C)'],
        marker='o', linestyle='-',label='Apparent Temperature (C)')
ax.plot(df1.loc['2006-04-01':'2016-04-01', 'Humidity'], marker='o',
        linestyle='-',label='Humidity')
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
ax.legend(loc = 'center right')
ax.set_xlabel('Month of April')
```

```
[11]: Text(0.5, 0, 'Month of April')
```




```
[12]: """We performed more closer analysis on the month april. We can see that the
      ↳apparent temperature does not fluctuate much
      and changes a little in the range 10-15(C). So from the plot there the
      ↳temperature is moderate without any dominant
      changes and the humidity remains the same like the rest of the year"""
      #Plotting each years Humidity and Temperature change
```

```
import matplotlib.dates as mdates
from datetime import datetime
```




```
fig, ax = plt.subplots(figsize=(15,5))
ax.plot(df1.loc['2006-04-01':'2016-04-01', 'Apparent Temperature (C)'],
        ↪marker='o', linestyle='-',label='Apparent Temperature (C)')
ax.plot(df1.loc['2006-04-01':'2016-04-01', 'Humidity'], marker='o',
        ↪linestyle='-',label='Humidity')
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
ax.legend(loc = 'center right')
ax.set_xlabel('Month of April')
```

[12]: Text(0.5, 0, 'Month of April')



```
[13]: sns.relplot(data = dat1, x = "Apparent Temperature (C)", y = "Humidity", color_
        ↪= 'purple')
```

[13]: <seaborn.axisgrid.FacetGrid at 0x28239bad648>



```
[14]: """From the rel plot we can see the intensity of humidity with Apparent_
      ↪temperature. We can see some instances of zero
      level humidity at very cold temperatures"""

dat2 = dat1.drop(['Formatted Date', 'Loud Cover', 'Daily Summary'], axis =1)
dat2.head()
```

```
[14]:
```

	Summary	Precip	Type	Temperature (C)	Apparent Temperature (C)	\
0	Partly Cloudy		rain	9.472222	7.388889	
1	Partly Cloudy		rain	9.355556	7.227778	
2	Mostly Cloudy		rain	9.377778	9.377778	
3	Partly Cloudy		rain	8.288889	5.944444	
4	Mostly Cloudy		rain	8.755556	6.977778	

	Humidity	Wind Speed (km/h)	Wind Bearing (degrees)	Visibility (km)	\
0	0.89	14.1197	251.0	15.8263	
1	0.86	14.2646	259.0	15.8263	
2	0.89	3.9284	204.0	14.9569	

3	0.83	14.1036	269.0	15.8263
4	0.83	11.0446	259.0	15.8263

Pressure (millibars)	
0	1015.13
1	1015.63
2	1015.94
3	1016.41
4	1016.51

[15]: *#in order to convert the categorical to dummy variables.*

```
dat2_dummy = pd.get_dummies(dat2)
dat2_dummy.head()
```

[15]:

	Temperature (C)	Apparent Temperature (C)	Humidity	Wind Speed (km/h)	\
0	9.472222	7.388889	0.89	14.1197	
1	9.355556	7.227778	0.86	14.2646	
2	9.377778	9.377778	0.89	3.9284	
3	8.288889	5.944444	0.83	14.1036	
4	8.755556	6.977778	0.83	11.0446	

	Wind Bearing (degrees)	Visibility (km)	Pressure (millibars)	\
0	251.0	15.8263	1015.13	
1	259.0	15.8263	1015.63	
2	204.0	14.9569	1015.94	
3	269.0	15.8263	1016.41	
4	259.0	15.8263	1016.51	

	Summary_Breezy	Summary_Breezy and Dry	Summary_Breezy and Foggy	...	\
0	0	0	0	...	
1	0	0	0	...	
2	0	0	0	...	
3	0	0	0	...	
4	0	0	0	...	

	Summary_Partly Cloudy	Summary_Rain	Summary_Windy	Summary_Windy and Dry	\
0	1	0	0	0	
1	1	0	0	0	
2	0	0	0	0	
3	1	0	0	0	
4	0	0	0	0	

	Summary_Windy and Foggy	Summary_Windy and Mostly Cloudy	\
0	0	0	
1	0	0	
2	0	0	

3	0	0
4	0	0


	Summary_Windy and Overcast	Summary_Windy and Partly Cloudy \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	Precip Type_rain	Precip Type_snow
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

[5 rows x 36 columns]

```
[61]: """Lets plot the correlation plot to determine the correlation between the_
↪variables"""
plt.figure(figsize=(12,8))
sns.heatmap(dat2_dummy.corr())
```

[61]: <matplotlib.axes._subplots.AxesSubplot at 0x2820e706208>



```
[ ]: """From the heatmap of the correlation plot we can see that Apparent_
    ↳Temperature is very positive correlated with the
Apparent Temperature denoting that they must have linear correlation between_
    ↳them, whereasHumidity is
very negatively correlated with the Apparent Temperature. Also we can see that_
    ↳visibility is negatively correlated
with Humidity.
```

```
[37]: #Now for the model

Y=dat2_dummy[['Apparent Temperature (C)']]
X=dat2_dummy.drop('Apparent Temperature (C)',axis=1)
```

```
[38]: Y.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 95936 entries, 0 to 96452
Data columns (total 1 columns):
Apparent Temperature (C)    95936 non-null float64
dtypes: float64(1)
```

memory usage: 1.5 MB

```
[ ]: sns.pairplot(dat2_dummy)
```

```
[39]: #Significance test
import statsmodels.api as sm
from statsmodels.stats import diagnostic as diag
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[40]: # define our input
X1 = sm.add_constant(X)
# create a OLS model
model = sm.OLS(Y, X1)
# fit the data
est = model.fit()
print(est.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:      Apparent Temperature (C)    R-squared:
0.990
Model:                                OLS    Adj. R-squared:
0.990
Method:                    Least Squares    F-statistic:
2.928e+05
Date:                    Tue, 08 Dec 2020    Prob (F-statistic):
0.00
Time:                    09:58:31    Log-Likelihood:
-1.4195e+05
No. Observations:                95936    AIC:
2.840e+05
Df Residuals:                95902    BIC:
2.843e+05
Df Model:                        33
Covariance Type:                nonrobust
=====
=====
```

			coef	std err	t
P> t	[0.025	0.975]			

const			-1.2965	0.058	-22.345
0.000	-1.410	-1.183			
Temperature (C)			1.1161	0.001	1933.943
0.000	1.115	1.117			
Humidity			0.7866	0.026	30.812
0.000	0.737	0.837			

Wind Speed (km/h)	-0.1045	0.001	-172.107
0.000 -0.106 -0.103			
Wind Bearing (degrees)	0.0005	3.23e-05	15.890
0.000 0.000 0.001			
Visibility (km)	0.0042	0.001	3.848
0.000 0.002 0.006			
Pressure (millibars)	0.0001	3.01e-05	4.961
0.000 9.02e-05 0.000			
Summary_Breezy	-1.0207	0.159	-6.403
0.000 -1.333 -0.708			
Summary_Breezy and Dry	2.2819	1.026	2.223
0.026 0.270 4.294			
Summary_Breezy and Foggy	-2.2499	0.190	-11.859
0.000 -2.622 -1.878			
Summary_Breezy and Mostly Cloudy	0.7164	0.089	8.041
0.000 0.542 0.891			
Summary_Breezy and Overcast	-0.2753	0.089	-3.097
0.002 -0.450 -0.101			
Summary_Breezy and Partly Cloudy	0.5212	0.093	5.613
0.000 0.339 0.703			
Summary_Clear	-0.4970	0.078	-6.383
0.000 -0.650 -0.344			
Summary_Dangerously Windy and Partly Cloudy	1.1514	1.027	1.121
0.262 -0.861 3.164			
Summary_Drizzle	-0.2496	0.181	-1.378
0.168 -0.605 0.105			
Summary_Dry	-1.4202	0.192	-7.402
0.000 -1.796 -1.044			
Summary_Dry and Mostly Cloudy	-1.3546	0.284	-4.764
0.000 -1.912 -0.797			
Summary_Dry and Partly Cloudy	-1.2062	0.135	-8.941
0.000 -1.471 -0.942			
Summary_Foggy	-0.1764	0.079	-2.244
0.025 -0.330 -0.022			
Summary_Humid and Mostly Cloudy	-0.6094	0.179	-3.397
0.001 -0.961 -0.258			
Summary_Humid and Overcast	-0.7111	0.395	-1.802
0.071 -1.484 0.062			
Summary_Humid and Partly Cloudy	-0.7095	0.260	-2.729
0.006 -1.219 -0.200			
Summary_Light Rain	-0.3391	0.150	-2.258
0.024 -0.633 -0.045			
Summary_Mostly Cloudy	-0.3960	0.077	-5.126
0.000 -0.547 -0.245			
Summary_Overcast	-0.4943	0.077	-6.387
0.000 -0.646 -0.343			
Summary_Partly Cloudy	-0.4978	0.077	-6.437
0.000 -0.649 -0.346			

Summary_Rain			-0.3451	0.333	-1.037
0.300	-0.998	0.307			
Summary_Windy			-0.1372	0.371	-0.370
0.711	-0.864	0.590			
Summary_Windy and Dry			1.4297	1.027	1.393
0.164	-0.582	3.442			
Summary_Windy and Foggy			2.0167	0.518	3.895
0.000	1.002	3.032			
Summary_Windy and Mostly Cloudy			1.9733	0.190	10.410
0.000	1.602	2.345			
Summary_Windy and Overcast			0.2519	0.171	1.472
0.141	-0.084	0.587			
Summary_Windy and Partly Cloudy			1.0507	0.147	7.140
0.000	0.762	1.339			
Precip Type_rain			-0.4263	0.030	-14.136
0.000	-0.485	-0.367			
Precip Type_snow			-0.8701	0.029	-29.515
0.000	-0.928	-0.812			

Omnibus:	2475.274	Durbin-Watson:	0.453
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2713.959
Skew:	0.384	Prob(JB):	0.00
Kurtosis:	3.296	Cond. No.	2.49e+16

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.64e-22. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[41]: import statsmodels.api as sm
cols = list(X.columns)
pmax = 1
while (len(cols)>0):
    p= []
    X_1 = X[cols]
    X_1 = sm.add_constant(X_1)
    model = sm.OLS(Y,X_1).fit()
    p = pd.Series(model.pvalues.values[1:],index = cols)
    pmax = max(p)
    feature_with_p_max = p.idxmax()
    if(pmax>0.05):
        cols.remove(feature_with_p_max)
    else:
        break
selected_features_BE = cols
```



```
print(selected_features_BE)
```

```
['Temperature (C)', 'Humidity', 'Wind Speed (km/h)', 'Wind Bearing (degrees)',
'Visibility (km)', 'Pressure (millibars)', 'Summary_Breezy', 'Summary_Breezy and
Dry', 'Summary_Breezy and Foggy', 'Summary_Breezy and Mostly Cloudy',
'Summary_Breezy and Partly Cloudy', 'Summary_Clear', 'Summary_Dry', 'Summary_Dry
and Mostly Cloudy', 'Summary_Dry and Partly Cloudy', 'Summary_Humid and Mostly
Cloudy', 'Summary_Humid and Partly Cloudy', 'Summary_Mostly Cloudy',
'Summary_Overcast', 'Summary_Partly Cloudy', 'Summary_Windy and Foggy',
'Summary_Windy and Mostly Cloudy', 'Summary_Windy and Overcast', 'Summary_Windy
and Partly Cloudy', 'Precip Type_rain', 'Precip Type_snow']
```

```
[42]: """As we can see that the Summary_Windy contains the maximum p-value and is
↳greater then the significance factor 0.05,
this variable is insignificant. So we need to remove summary_windy count and
↳keep checking the p-values by removing
insignificant variables until p-value is closer to zero"""
X1_opt = X1.drop('Summary_Windy', axis=1)
X1_opt=X1_opt.drop('Summary_Foggy',axis=1)
X1_opt=X1_opt.drop('Summary_Rain',axis=1)
X1_opt=X1_opt.drop('Summary_Drizzle',axis=1)
X1_opt=X1_opt.drop('Summary_Dangerously Windy and Partly Cloudy',axis=1)
X1_opt=X1_opt.drop('Summary_Light Rain',axis=1)
X1_opt=X1_opt.drop('Summary_Windy and Dry',axis=1)
X1_opt=X1_opt.drop('Summary_Humid and Overcast',axis=1)

# create a OLS model
model = sm.OLS(Y, X1_opt)
# fit the data
est = model.fit()
print(est.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:    Apparent Temperature (C)    R-squared:
0.990
Model:                                OLS    Adj. R-squared:
0.990
Method:                                Least Squares    F-statistic:
3.716e+05
Date:                                Tue, 08 Dec 2020    Prob (F-statistic):
0.00
Time:                                09:58:43    Log-Likelihood:
-1.4195e+05
No. Observations:                                95936    AIC:
```

2.840e+05

Df Residuals: 95909 BIC:

2.842e+05

Df Model: 26

Covariance Type: nonrobust

[0.025 0.975]		coef	std err	t	P> t

const		-1.4147	0.028	-50.154	0.000
-1.470	-1.359				
Temperature (C)		1.1160	0.001	1934.439	0.000
1.115	1.117				
Humidity		0.7856	0.026	30.778	0.000
0.736	0.836				
Wind Speed (km/h)		-0.1044	0.001	-172.424	0.000
-0.106	-0.103				
Wind Bearing (degrees)		0.0005	3.23e-05	15.897	0.000
0.000	0.001				
Visibility (km)		0.0042	0.001	3.800	0.000
0.002	0.006				
Pressure (millibars)		0.0001	3e-05	4.968	0.000
9.02e-05	0.000				
Summary_Breezy		-0.8422	0.147	-5.743	0.000
-1.130	-0.555				
Summary_Breezy and Dry		2.4605	1.063	2.315	0.021
0.377	4.544				
Summary_Breezy and Foggy		-2.0718	0.181	-11.457	0.000
-2.426	-1.717				
Summary_Breezy and Mostly Cloudy		0.8952	0.052	17.344	0.000
0.794	0.996				
Summary_Breezy and Overcast		-0.0964	0.051	-1.891	0.059
-0.196	0.003				
Summary_Breezy and Partly Cloudy		0.7000	0.058	11.984	0.000
0.585	0.814				
Summary_Clear		-0.3178	0.020	-16.164	0.000
-0.356	-0.279				
Summary_Dry		-1.2413	0.183	-6.768	0.000
-1.601	-0.882				
Summary_Dry and Mostly Cloudy		-1.1758	0.285	-4.129	0.000
-1.734	-0.618				
Summary_Dry and Partly Cloudy		-1.0273	0.116	-8.825	0.000
-1.255	-0.799				
Summary_Humid and Mostly Cloudy		-0.4299	0.169	-2.546	0.011
-0.761	-0.099				
Summary_Humid and Partly Cloudy		-0.5300	0.258	-2.052	0.040

-1.036	-0.024				
Summary_Mostly Cloudy		-0.2169	0.018	-12.302	0.000
-0.251	-0.182				
Summary_Overcast		-0.3153	0.017	-18.289	0.000
-0.349	-0.281				
Summary_Partly Cloudy		-0.3186	0.018	-17.534	0.000
-0.354	-0.283				
Summary_Windy and Foggy		2.1951	0.532	4.126	0.000
1.152	3.238				
Summary_Windy and Mostly Cloudy		2.1520	0.181	11.859	0.000
1.796	2.508				
Summary_Windy and Overcast		0.4305	0.160	2.683	0.007
0.116	0.745				
Summary_Windy and Partly Cloudy		1.2293	0.133	9.274	0.000
0.969	1.489				
Precip Type_rain		-0.4856	0.016	-29.517	0.000
-0.518	-0.453				
Precip Type_snow		-0.9291	0.015	-62.269	0.000
-0.958	-0.900				

Omnibus:	2473.706	Durbin-Watson:	0.453
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2711.911
Skew:	0.384	Prob(JB):	0.00
Kurtosis:	3.295	Cond. No.	9.33e+16

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.16e-23. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[43]: """Removing the constant term"""
```

```
X1_opt2 = X1_opt.drop('const', axis=1)
```

```
[44]: """Linear Regression model"""
```

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

X_train, X_test, Y_train, Y_test = train_test_split(X1_opt2, Y, test_size = 0.
↳5, random_state = 1)
```

```
[45]: print(X_train.shape)
```

```
print(Y_train.shape)
```

```
(47968, 27)
```

```
(47968, 1)
```

```
[46]: lm = LinearRegression()
lm.fit(X_train, Y_train)
```

```
[46]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[47]: Y_pred = lm.predict(X_test)
```

```
[48]: Y_pred_dataframe = pd.DataFrame(Y_pred, columns = ['Predicted_App Temp'])
Y_pred_dataframe.head()
```

```
[48]: Predicted_App Temp
0      6.969721
1      5.258607
2     15.935838
3     23.027109
4      2.267791
```

```
[49]: Y_test
```

```
[49]: Apparent Temperature (C)
37747      6.405556
50630      6.161111
91632     15.444444
30623     22.394444
71677      1.838889
...
15415     -1.816667
38415     -8.488889
28768     -5.794444
45247     19.933333
27644     18.088889
```

```
[47968 rows x 1 columns]
```

```
[50]: #lets grab the coefficients and intercept
Y_coeff = lm.coef_
Y_intercept = lm.intercept_
print(Y_coeff)
```

```
[[ 1.11573740e+00  7.35990645e-01 -1.04344707e-01  4.99374599e-04
  4.31574116e-03  1.69615367e-04 -1.11342986e+00  2.43973205e+00
 -1.83824067e+00  9.43096359e-01 -5.62890487e-02  6.78884920e-01
 -3.00063866e-01 -1.41017203e+00 -1.07339950e+00 -1.15543282e+00
 -5.63812735e-01 -4.80990781e-01 -2.12534186e-01 -2.99775469e-01
 -3.15938611e-01  1.81432460e+00  2.26193720e+00  6.65760108e-01
  1.02846618e+00  2.16353837e-01 -2.16353837e-01]]
```

```
[51]: from sklearn.metrics import mean_squared_error, r2_score
print(Y_intercept)
r2 = r2_score(Y_test, Y_pred)
print(lm.score(X_test, Y_test))
print("R^2 score: ", r2)
print('RMSE:', np.sqrt(mean_squared_error(Y_test, lm.predict(X_test))))
```


[-2.09910272]

0.9901798864537803

R^2 score: 0.9901798864537803

RMSE: 1.0614980989251768

```
[52]: #DataFrame.as_matrix(columns=None)
Y_test.head()
plt.scatter(X_test['Temperature (C)'].values, Y_test.values, color = "blue" )
plt.plot(X_test['Temperature (C)'].values, Y_pred, color = "red")
plt.title("Truth or bluff (Linear Regression)")
plt.xlabel('Temperature')
plt.ylabel('Apparent Temperature')
plt.show()
```



```
[53]: Y_pred
```

```
[53]: array([[ 6.96972075],
           [ 5.25860737],
           [15.93583807],
           ...,
           [-4.49590698],
           [19.67576139],
           [17.97798901]])
```

```
[54]: """First we will create a matrix of different containing nonlinear features"""
from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 2)
"""x1 as 1st feature, x2 as second feature"""
X_poly = poly_reg.fit_transform(X_train)#b0+b1x1+b2x2^2
lm_poly = LinearRegression()
lm_poly.fit(X_poly,Y_train)
```

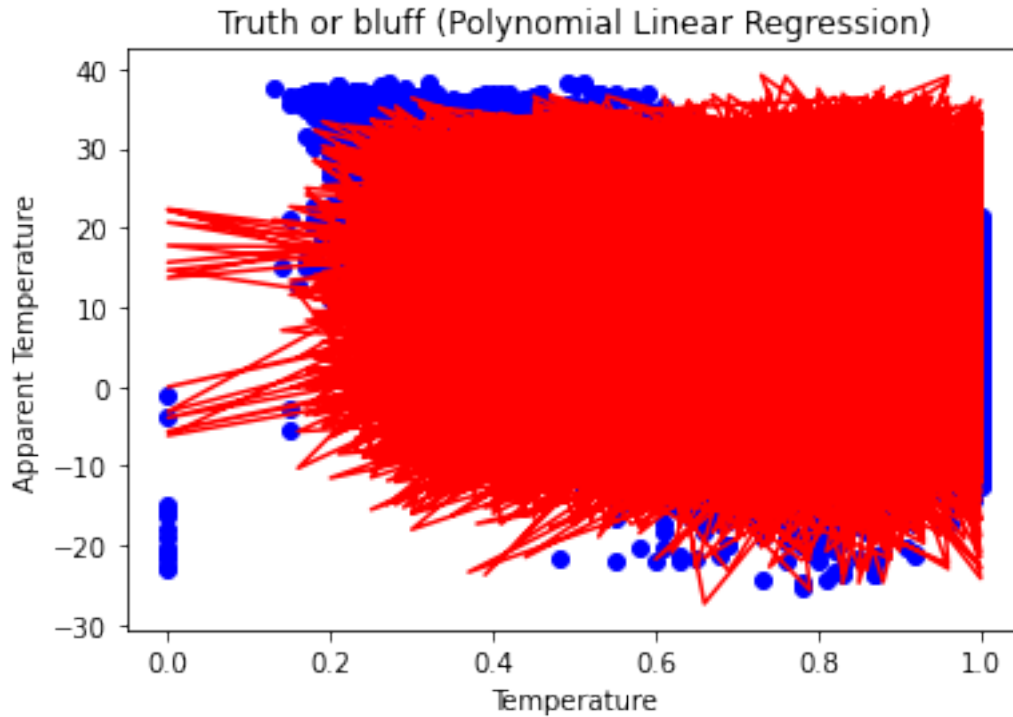
```
[54]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[55]: X_train['Humidity']
```

```
[55]: 71989    0.65
39793    0.93
12157    0.67
82488    0.72
26755    0.33
...
21440    0.67
73492    0.92
50057    0.65
5192     0.89
77851    0.96
Name: Humidity, Length: 47968, dtype: float64
```

```
[56]: Y_poly_pred = lm_poly.predict(X_poly)
```

```
[57]: plt.scatter(X_test['Humidity'],Y_test, color = "blue" )
plt.plot(X_test['Humidity'],Y_poly_pred, color = "red")
plt.title("Truth or bluff (Polynomial Linear Regression)")
plt.xlabel('Temperature')
plt.ylabel('Apparent Temperature')
plt.show()
```



```
[69]: from sklearn.metrics import mean_squared_error, r2_score
r2 = -r2_score(Y_test, Y_poly_pred)
#print(lm_poly.score(X_test, Y_test))
print('R^2 score:', r2)
print('RMSE:', np.sqrt(mean_squared_error(Y_test, Y_poly_pred)))
```

R² score: 0.9911999745738618

RMSE: 15.11535686350149

```
[ ]: """We can see that the polynomial regression performs a little bit better as
↳ the R^2 score is 0.99017 and 0.9911 for
multiple regression and polynomial regression respectively"""
```

```
[ ]: """Now we will apply ridge regression and lasso model to the dataset"""
```

```
[70]: from sklearn.linear_model import Ridge
from sklearn import model_selection
from sklearn.linear_model import RidgeCV
```

```
[71]: lm_ridge = Ridge()
lm_ridge.fit(X_train, Y_train)
```

```
[71]: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,  
          normalize=False, random_state=None, solver='auto', tol=0.001)
```

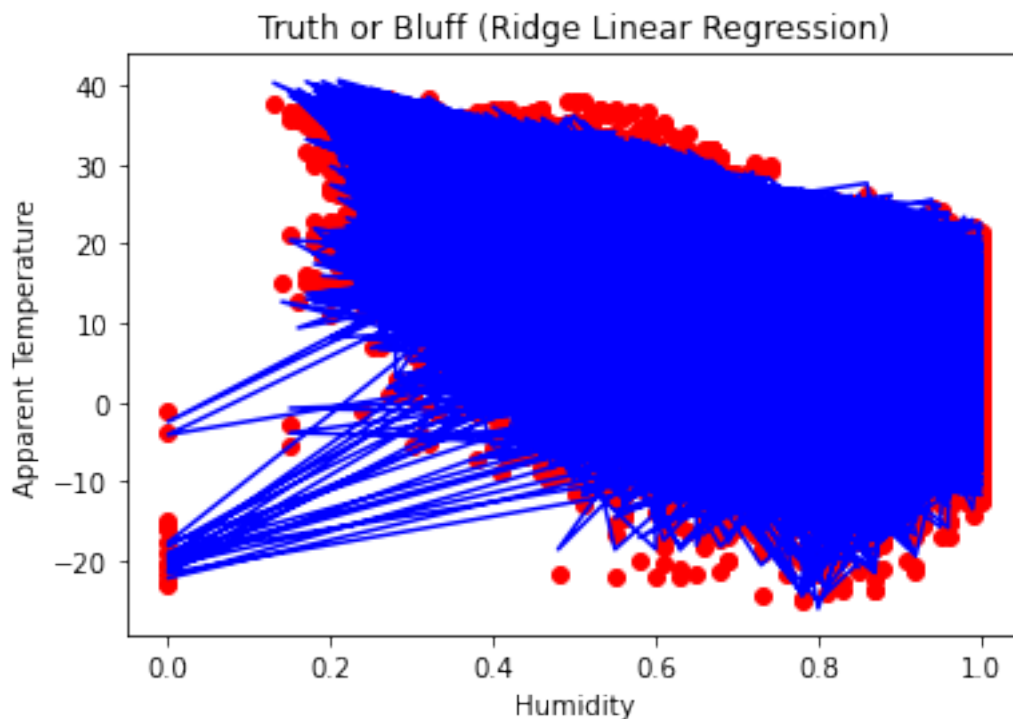
```
[72]: lm_ride.coef_
```

```
[72]: array([[ 1.11572046e+00,  7.36007021e-01, -1.04253523e-01,  
            4.99560931e-04,  4.31536851e-03,  1.70605935e-04,  
           -1.07803452e+00,  1.21881135e+00, -1.73216556e+00,  
            9.37343303e-01, -5.82897501e-02,  6.73203411e-01,  
           -2.99812778e-01, -1.34325586e+00, -9.20302546e-01,  
           -1.13089373e+00, -5.33961094e-01, -4.43956363e-01,  
           -2.12692300e-01, -3.00016258e-01, -3.15922070e-01,  
            1.35833272e+00,  2.07064198e+00,  6.37045639e-01,  
            9.89884640e-01,  2.16528251e-01, -2.16528251e-01]])
```

```
[73]: lm_ride.intercept_
```

```
[73]: array([-2.10095029])
```

```
[74]: plt.scatter(X_test['Humidity'], Y_test, color = 'red')  
plt.plot(X_test['Humidity'], lm_ride.predict(X_test), color = 'blue')  
plt.title('Truth or Bluff (Ridge Linear Regression)')  
plt.xlabel('Humidity')  
plt.ylabel('Apparent Temperature')  
plt.show()
```




```
[77]: """We can see that most of the predicted values well converge with the test_
↪values"""
      """Lets compute the R^2 score for validity"""

MSE_Ridge = metrics.mean_squared_error(Y_test, lm_ridge.predict(X_test))
print('MSE:', MSE_Ridge)
print('RMSE:', np.sqrt(metrics.mean_squared_error(Y_test, lm_ridge.
predict(X_test))))
r2 = r2_score(Y_test, lm_ridge.predict(X_test))
#print(lm_poly.score(X_test, Y_test))
print('R^2 score:', r2)
```

```
MSE: 1.126815266682861
RMSE: 1.061515551785682
R^2 score: 0.9901795635319004
```

```
[80]: """We can see that R^2 score is 0.99018 which is denotes that ridge model_
↪performs very well"""
      """We'll use cross validation to determine the optimal alpha value. By default,
↪the ridge regression
cross validation class uses the Leave One Out strategy (k-fold).
We can compare the performance of our model with different alpha
values by taking a look at the mean square error."""
from sklearn.linear_model import RidgeCV
alphas = 10*np.linspace(10,-2,100)*0.5
#alphas=[0, 1, 1000, 1000000]
lm_ridge_opt = RidgeCV(alphas=alphas,scoring = "neg_mean_squared_error", cv_
↪=10, normalize = True)
lm_ridge_opt.fit(X_train, Y_train)
print("optimum shrinkage parameter using Ridge CV",lm_ridge_opt.alpha_)
print("MSE for Ridge CV",metrics.mean_squared_error(Y_test, lm_ridge_opt.
↪predict(X_test)))
print("RMSE for RidgeCV",np.sqrt(metrics.mean_squared_error(Y_test,
↪lm_ridge_opt.predict(X_test))))
```

```
optimum shrinkage parameter using Ridge CV 0.005
MSE for Ridge CV 1.1345524643911613
RMSE for RidgeCV 1.0615137280557963
```

```
[81]: r2 = r2_score(Y_test, lm_ridge_opt.predict(X_test))
#print(lm_poly.score(X_test, Y_test))
print('R^2 score:', r2)
```

```
R^2 score: 0.9901121321961862
```

```
[82]: """So we can see the shrinkage parameter using ridge CV is about 0.05, with
      ↪ this optimal value we get good the RMSE value. Hence we choose alpha = 0.005
      ↪ the
      optimal shrinkage parameter"""

"""Now we perform the Lasso"""
from sklearn.linear_model import Lasso, LassoCV
lassocv = LassoCV(alphas = alphas, cv = 10, max_iter = 100000, normalize = True)
lassocv.fit(X_train, Y_train)
#lasso.set_params(alpha=lassocv.alpha_)
#lasso.fit(X_opt, Y_opt)
print("MSE",metrics.mean_squared_error(Y_test, lasso.predict(X_test)))
print("RMSE",np.sqrt(metrics.mean_squared_error(Y_test, lasso.
      ↪ predict(X_test))))
print("optimal shrinkage parameter for Lasso Model:", lasso.alpha_)
```

MSE 2.8963302859859095
 RMSE 1.7018608303812357
 optimal shrinkage parameter for Lasso Model: 0.005

```
[83]: r2 = r2_score(Y_test,lassocv.predict(X_test))
      #print(lm_poly.score(X_test,Y_test))
      print('R^2 score:',r2)
```

R^2 score: 0.9747578610219854

```
[ ]: """For the lasso cv model we can see the optimal shrinkage parameter is same
      ↪ and it is 0.005"""

"""From the lasso model we can see that the R^2 score is 0.974, hence the
      ↪ performance decline for this model. And multiple
      regression , polynomial regression and ridge regression performs better then
      ↪ this model"""
```

```
[86]: """Now we will apply another model Random forest with maximum depth 10"""
from sklearn.ensemble import RandomForestRegressor
lm_rf = RandomForestRegressor(max_depth= 10, random_state = 0, n_estimators =
      ↪ 10)
lm_rf.fit(X_train, Y_train)
```

```
[86]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=10,
                             max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=10,
                             n_jobs=None, oob_score=False, random_state=0, verbose=0,
                             warm_start=False)
```

```
[87]: r2 = r2_score(Y_test,lm_rf.predict(X_test))
      #print(lm_poly.score(X_test,Y_test))
      print('R^2 score:',r2)
```

R^2 score: 0.9998736673437375

```
[89]: """Random Forest with max depth 50"""
      from sklearn.ensemble import RandomForestRegressor
      lm_rf_50 = RandomForestRegressor(max_depth= 50, random_state = 0, n_estimators=
      ↪ 10)
      lm_rf_50.fit(X_train, Y_train)
```

```
[89]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=50,
                             max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=10,
                             n_jobs=None, oob_score=False, random_state=0, verbose=0,
                             warm_start=False)
```

```
[90]: r2 = r2_score(Y_test,lm_rf_50.predict(X_test))
      #print(lm_poly.score(X_test,Y_test))
      print('R^2 score:',r2)
```

R^2 score: 0.9999459787231981

```
[93]: """Conclusion"""
      from astropy.table import QTable, Table, Column
      data_rows = [('Multiple Linear Regression', r2_score(Y_test,Y_pred)),(
      ↪ 'Polynomial Regression', -r2_score(Y_test,Y_poly_pred)),
                  ('Ridge Regression', r2_score(Y_test,lm_ridge_opt.
      ↪ predict(X_test))),
                  ('Lasso Model', r2_score(Y_test,lassocv.predict(X_test))),
                  ('Random Forest model_ with depth 10', r2_score(Y_test,lm_rf.
      ↪ predict(X_test))),
                  ('Random Forest model_ depth 50',r2_score(Y_test,lm_rf_50.
      ↪ predict(X_test))) ]
      t = Table(rows=data_rows, names=('Model', 'R^2_score'))
      print(t)
```

Model	R^2_score
Multiple Linear Regression	0.9901798864537803
Polynomial Regression	0.9911999745738618
Ridge Regression	0.9901121321961862
Lasso Model	0.9747578610219854
Random Forest model_ with depth 10	0.9998736673437375

Random Forest model_ depth 50 0.9999459787231981

```
[ ]: """This is the summary of the models used and  $R^2$  score. From the summary we  
→ can see that the Polynomial Linear  
regression performs well with comparison to Multiple linear and ridge  
→ regression for this dataset. But the Random  
Forest model with depth 50 performs the best with  $r^2$  score of 0.9999"""
```